

Decision Tree Classifier – Custom Implementation

Programming in Python Language
Michał Janik AGH

1. Project Overview

The goal of this project is to implement a simple, fully transparent **Decision Tree Classifier** using only Python and NumPy.

Unlike scikit-learn, this implementation is written **entirely from scratch**, focusing on full algorithmic understanding rather than performance.

The tree uses the **CART methodology**, based on binary splits and impurity minimization (Gini or Entropy).

It supports configurable parameters such as maximum depth and minimum samples required to split.

2. Project Structure

The project is split into modules for clarity and maintainability:

- **decision_tree.py**

Contains the full implementation of:

- impurity functions (gini, entropy),
- the Node dataclass used as the internal structure of the tree,
- DecisionTreeClassifier containing training, prediction, and scoring logic,
- optional comparison function using scikit-learn (compare).

- **tests/**

Includes unit and integration tests for:

- impurity correctness,
 - split logic,
 - end-to-end tree training,
 - compatibility tests with scikit-learn (if installed).
-

3. Algorithm Description

3.1. Dataset Overview

The implementation is dataset-agnostic.

It accepts any NumPy array x (shape: $n_{\text{samples}} \times n_{\text{features}}$) and integer labels array y .

Typical examples used for development and testing:

- simple binary classification (0/1),
- small synthetic datasets,
- Iris dataset (via scikit-learn, optional).

This keeps the algorithm simple and reproducible.

3.2. Data Preprocessing

The classifier assumes minimal preprocessing:

1. **Ensure that feature values in x are numeric**
(Decision Trees cannot process strings directly).
2. **Convert labels to integers**
(`fit()` handles this internally using `np.array(y, dtype=int)`).
3. **Optional normalization**
Not required, because decision trees are scale-invariant.

Unlike deep learning models, decision trees handle raw values without augmentation or normalization.

3.3. Tree Construction Algorithm

The decision tree implements the following components:

- **Impurity Metrics**

Two impurity functions are available:

- **Gini impurity**
- **Shannon entropy**

Both measure how mixed the labels are in a node.

- **Split Selection Strategy**

The tree performs **exhaustive search** across:

- every feature,
- every possible split point (midpoint between sorted unique values).

For each candidate split:

1. The dataset is divided into left/right nodes.
2. Both sides' impurity is calculated.
3. A weighted impurity score is computed.
4. Impurity reduction (gain) is measured.
5. The best split is selected.

This mimics the logic used by CART.

• **Stopping Conditions**

The recursion stops when:

- all labels in the node are identical,
- maximum depth is reached,
- no split yields a positive impurity gain,
- one side of the split would contain fewer samples than `min_samples_split`.

In these cases the node becomes a **leaf** with a majority label.

• **Tree Representation**

Each node is represented by:

```
@dataclass
class Node:
    feature
    threshold
    left
    right
    value # used only in leaf nodes
```

This makes the structure lightweight and easy to print or visualize.

3.4. Training Process

Training is handled by:

```
DecisionTreeClassifier.fit(X, y)
```

Key elements:

- Recursively builds the tree using `_build()`.

- Node purity and depth limits are respected.
- Majority voting is used for leaf assignment.
- The resulting tree is stored as `self.tree_`.

The training process is deterministic (no randomness).

3.5. Evaluation and Inference

Prediction works by traversing the tree:

- starting from the root,
- comparing a feature to its threshold,
- moving left or right until reaching a leaf node.

Accuracy is measured with:

`score(X, y)`

which computes standard classification accuracy.

3.6. Optional Comparison with Scikit-learn

If scikit-learn is installed, running:

`compare(max_depth=3)`

will:

1. Train this custom classifier on the Iris dataset.
2. Train scikit-learn's `DecisionTreeClassifier`.
3. Print both accuracies.

This allows validating correctness and experiment-level performance.

4. Installation

To run the project, set up a Python virtual environment and install dependencies.

```
python -m venv .venv
.\venv\Scripts\activate # Windows
# or
source .venv/bin/activate # Linux/Mac

pip install numpy scikit-learn pytest
```

5. How to Run the Code

Train and test the tree manually

```
python decision_tree.py
```

Run scikit-learn comparison

```
python decision_tree.py --demo
```

Run the unit tests

```
pytest tests/
```

The project can be extended by adding plotting, exporting the tree structure, or supporting additional impurity metrics.